


Current baselines across GPUs: Longer-MQAR accuracy (best.pt)



Mean ± population SD over training seeds 123/124/125 (n=3 per GPU); 500 examples per slice; FP32, TF32 off. Training and eval both ran on the labeled GPU.